

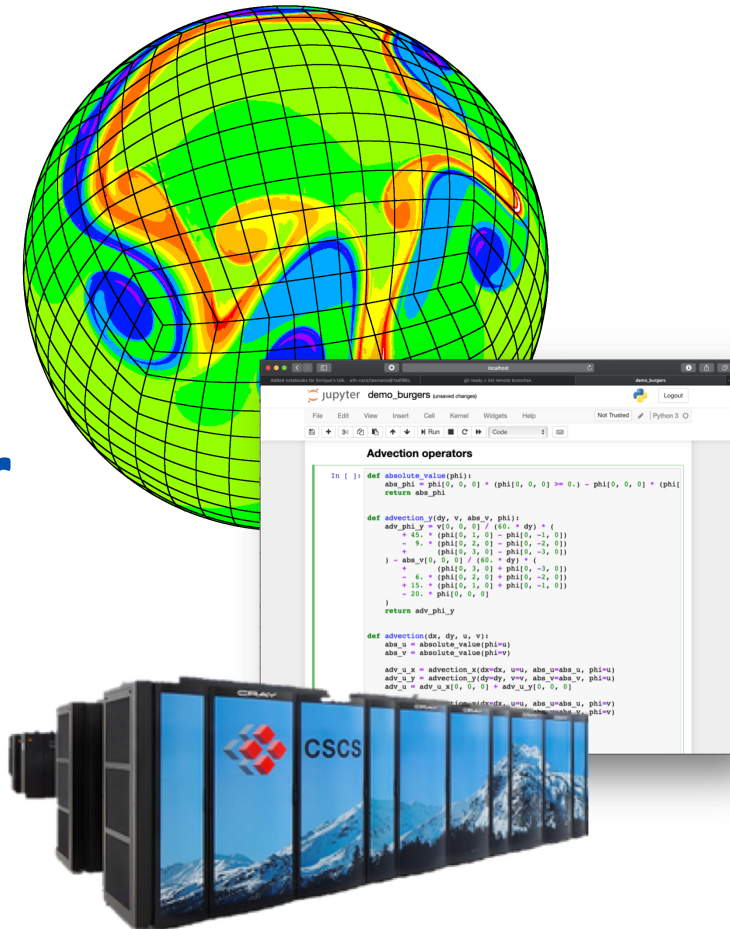
# High Performance Computing for Weather and Climate (HPC4WC)

Content: High-Level Programming

Lecturer: Stefano Ubbiali

Block course 701-1270-00L

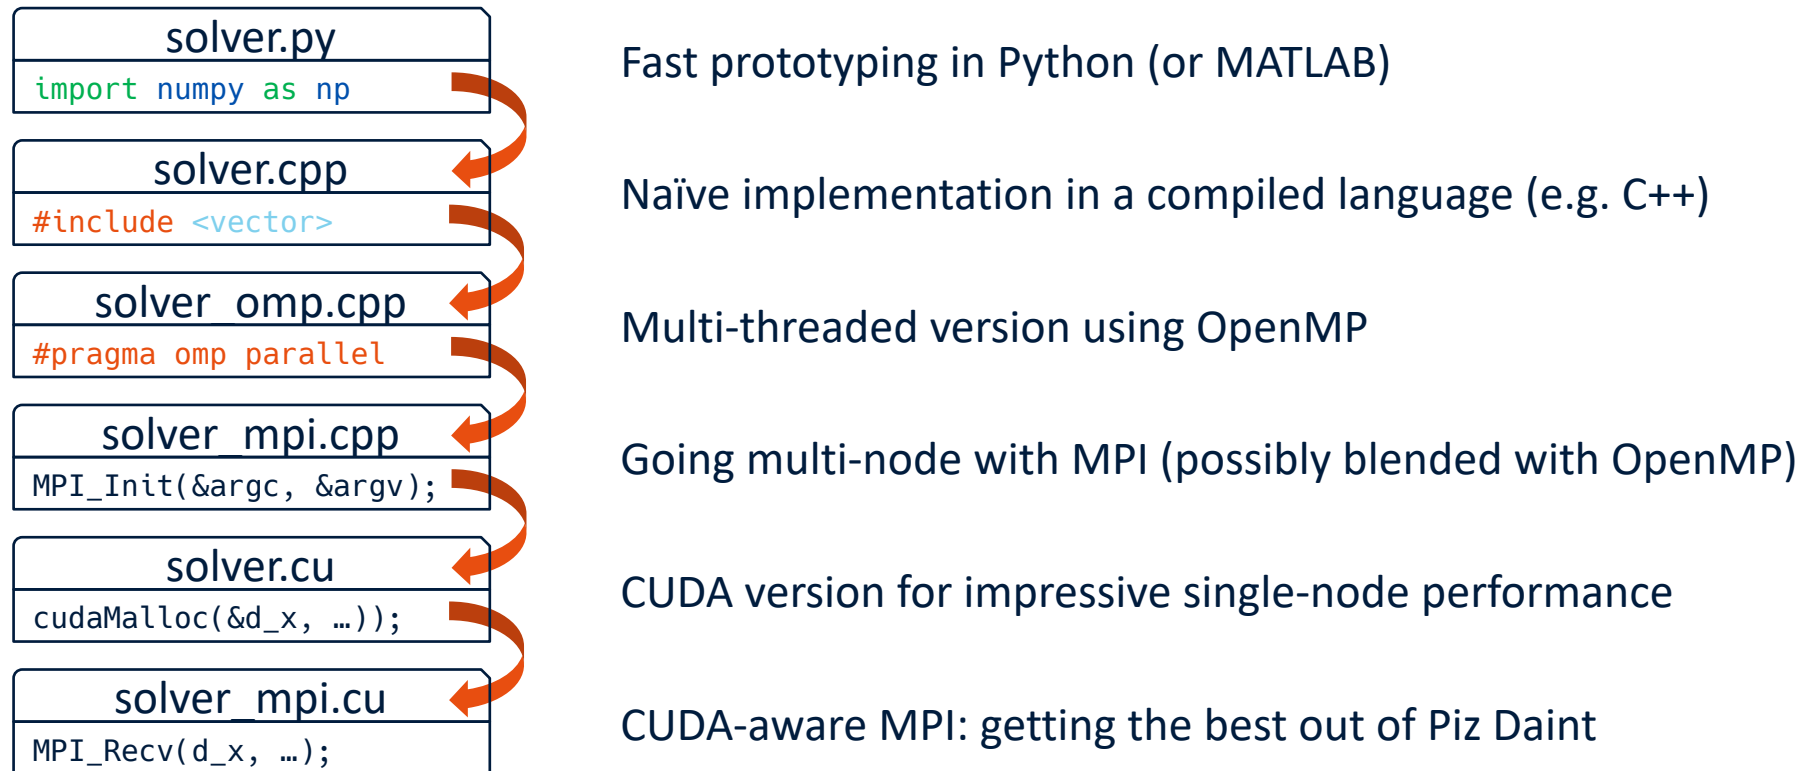
Summer 2020



# Learning Goals

- Learning what a domain-specific language (DSL) is.
- Understanding how a DSL helps in writing hardware-agnostic, maintainable code without sacrificing performance.
- Be able to apply a DSL to a stencil program from a weather and climate model.

# Typical Workflow



Why is this approach problematic?

# Possible Scenarios

What if ...

1. ... we want to introduce a modification at the algorithmic/numerical level?
2. ... our application has a broad user community and it must run efficiently on a variety of platforms?
3. ... our code consists of thousands (if not millions) LOC?

**The explosion of hardware architectures made this development model obsolete!**

# A Real-Case Example : COSMO

- Limited-area model developed by the **C**onsortium for **S**mall-Scale **M**odeling.
- Run operationally by 7 national weather services and used by several academic institutions as a research tool.
- Two target architectures: CPUs and GPUs.
- Around 330K lines of F90 code and 90K lines of C/C++ code.
- Cost of porting the full code base to GPU: approx. 20-30 Man-Years!

# Separation of Concerns

## Domain expert

Answer scientific research questions

Declarative programming style:  
Focus on **what** you want to do

Common data access interface:  
e.g. `data[i, j, k]`

Computation kernels:  
Calculations for a single grid point

Individual operators (“grains”)

## Performance expert

Write optimized code for target platform

Imperative programming style:  
Focus on **how** to do it

Storage and memory allocation:  
e.g. C-layout vs F-layout

Control structure (e.g. for loops):  
Optimized data traversal

Final computation:  
Detect and exploit parallelism b/w grains

# Overarching Goals

- Single **hardware-agnostic** application code.
- Easy to implement.
- Easy to **read**.
- Easy to **maintain**.
- **Performance portable**.



# Domain Specific Languages (DSLs)

- Programming language tailored for a specific class of problems.
- Higher level of abstraction w.r.t. a general purpose language.
- Intended to be used by domain experts, who may not be fluent in programming.
- Abstractions and notations much aligned to concepts and rules from the domain.
- Some examples:
  - Machine Learning: TensorFlow (Keras)
  - Scientific Computing: Kokkos, FEniCS, FreeFEM
  - Fluid Dynamics: OpenFOAM
  - Image Processing: Halide
  - Stencils: Ebb, Taichi, GT4Py

# GT4Py

- High-performance implementation of a stencil kernel from a high-level definition.
- GT4Py is a domain specific **library** which exposes a domain specific **language** (GTScript) to express the stencil logic.
- GTScript is embedded in Python (**eDSL**).
  - Legal Python syntax and (almost) legal Python semantics.
- GT4Py = **GridTools For Python**
  - Harnessing the C++ GridTools ecosystem to generate native implementations of the stencils.
- Emphasis on tight integration with scientific Python stack.

# Definitions Function

```
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
import numpy as np
```

```
f64 = np.float64
```

```
def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0,  0, 0]
            +   in_field[-1,  0, 0]
            +   in_field[+1,  0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Regular (named) function

# Definitions Function

```
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
import numpy as np

f64 = np.float64

def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0,  0, 0]
            +   in_field[-1,  0, 0]
            +   in_field[+1,  0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Input and output fields  
(object-oriented interface)

# Definitions Function

```
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
import numpy as np

f64 = np.float64

def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0, 0, 0]
            +   in_field[-1, 0, 0]
            +   in_field[+1, 0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Field descriptors  
as type annotations

# Definitions Function

```
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
import numpy as np

f64 = np.float64

def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0,  0, 0]
            +   in_field[-1,  0, 0]
            +   in_field[+1,  0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Any computation must be wrapped in a **with** construct which can be thought of as being a k-loop

# Definitions Function

```
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
import numpy as np

f64 = np.float64

def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0,  0, 0]
            +   in_field[-1,  0, 0]
            +   in_field[+1,  0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Iteration order in the  
vertical direction :  
PARALLEL, FORWARD,  
BACKWARD

# Definitions Function

```
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
import numpy as np

f64 = np.float64

def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0, 0, 0]
            +   in_field[-1, 0, 0]
            +   in_field[+1, 0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Vertical region of application:  
... = full column



# Definitions Function

```
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
import numpy as np
```

```
f64 = np.float64
```

```
def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
```

```
        out_field = (
            - 4. * in_field[ 0,  0, 0]
            +   in_field[-1,  0, 0]
            +   in_field[+1,  0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Each statement (or **stage**) can  
be thought of as an ij-loop

# Definitions Function

```
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
import numpy as np

f64 = np.float64

def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0, 0, 0]
            +   in_field[-1, 0, 0]
            +   in_field[+1, 0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

Neighboring points accessed  
through offsets

# Definitions Function

```
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
import numpy as np

f64 = np.float64

def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[0, 0, 0]
            + in_field[-1, 0, 0]
            + in_field[+1, 0, 0]
            + in_field[ 0, -1, 0]
            + in_field[ 0, +1, 0] )
```

1<sup>st</sup> horizontal dimension

2<sup>nd</sup> horizontal dimension

Vertical dimension

Neighboring points accessed through offsets

# Definitions Function

```
import gt4py as gt
from gt4py.gtscript import Field, PARALLEL, computation, interval
import numpy as np

f64 = np.float64

def laplacian_defs(in_field: Field[f64], out_field: Field[f64]):
    with computation(PARALLEL), interval(...):
        out_field = (
            - 4. * in_field[ 0,  0, 0]
            +   in_field[-1,  0, 0]
            +   in_field[+1,  0, 0]
            +   in_field[ 0, -1, 0]
            +   in_field[ 0, +1, 0] )
```

No for loops!  
No return statement!

# Compilation

- A stencil needs to be compiled for a given **backend**:

```
backend = "gtx86"  
laplacian = gt.stencil(backend, laplacian_defs)
```

- Available backends:
  - Python: "debug" (for loops), "numpy" (vectorized syntax);
  - C++: "gtx86" (x86), "gtmc" (MIC), "gtcuda" (NVIDIA GPU).
- For GT-based backends, compilation consists of three steps:
  - 1) Generate optimized code for the target architecture.
  - 2) Compile the automatically generated code.
  - 3) Build Python bindings to that code.

# Storages

- The compilation returns a callable object which can be invoked on GT4Py storages.
- Storages have optimal memory **strides**, **alignment** and **padding**.
- `gt.storage` provides functionalities to allocate storages ...

```
nx, ny, nz = 128, 128, 64
def_orig = (1, 1, 0)
out_field = gt.storage.zeros(
    backend, def_orig, (nx, ny, nz), dtype=f64 )
```

... and convert NumPy arrays into valid storages:

```
in_field = gt.storage.from_array(
    np.random.rand(nx, ny, nz), backend, def_orig, dtype=f64 )
```

# Storages

- Storages can be accessed as NumPy arrays:

```
in_field[0, 0, 0] = 4.  
print(in_field[0, 0, 0])
```

# Running

- Running computations is as simple as a function call:

```
laplacian(  
    in_field=in_field,  
    out_field=out_field,  
    origin=(1, 1, 0),  
    domain=(nx - 2, ny - 2, nz)  
)
```

Bindings b/w the symbols used  
within the definitions fct.  
and the arrays holding the data



# Running


- Running computations is as simple as a function call:

```
laplacian(  
    in_field=in_field,  
    out_field=out_field,  
    origin=(1, 1, 0),  
    domain=(nx - 2, ny - 2, nz)  
)
```


Origin and extent of the  
computation domain

- `out_field` now contains the results of the computation.


# Region of application




# Region of application




# Region of application




# Region of application




# Region of application




# Region of application




# Region of application





 (Largest) computation domain




# Region of application




-  (Largest) computation domain
-  (Thinnest) boundary region (halo)

# Region of application



# Region of application



# Disadvantages of a DSL

- Lack of generality: A DSL is not a complete ontology!
- Debugging on the generated code.
- Cost of developing and maintaining the DSL compiler toolchain.

# Conclusions

- High-level programming techniques hide the complexities of the underlying architecture to the end user.
- DSL allows to target multiple platforms without polluting the application code with hardware-specific boilerplate code.
- GT4Py is a Python framework to write performance portable applications in the weather and climate area. It ships with a DSL to write stencil computations.

# Lab Exercises

## 01-GT4Py-sumdiff.ipynb

- Compare NumPy, CuPy and GT4Py on the sum-diff stencil (demo).

## 02-GT4Py-laplacian.ipynb

- Compare NumPy, CuPy and GT4Py on the Laplacian stencil (demo).

## 03-GT4Py-concepts.ipynb

- Digest the main concepts of GT4Py.
- Get familiar with writing, compiling and running stencils.
- Get insights on the internal data-layout of the storages.

## 04-GT4Py-stencil2d.ipynb

- Step-by-step porting of stencil2d.py to GT4Py.
- Write two alternative versions of stencil2d-gt4py-v0.py

# Before Starting

1. Pull the latest commit from the Github repo.
2. Make sure that your `.jupyterhub.env` contains the following lines:

```
module load Boost
module load cudatoolkit
NVCC_PATH=$(which nvcc)
CUDA_PATH=$(echo $NVCC_PATH | sed -e "s/\\/bin\\/nvcc//g")
export CUDA_HOME=$CUDA_PATH
export LD_LIBRARY_PATH=$CUDA_PATH/lib64:$LD_LIBRARY_PATH
```

After updating your `.jupyterhub.env` from a terminal (see Oli's post in #general):  
terminate the JupyterLab session and fire up a new one.

Have fun with GT4Py!



# References

Broad introduction to DSLs:

<https://www.jetbrains.com/mps/concepts/domain-specific-languages/>

Designing APIs - The Case of GridTools (M. Bianco):

[https://www.youtube.com/watch?v=IzWxgFcJFdk&list=PL1tk5lGm7zvQOXi24s586pwDF\\_yseZ-80&index=7](https://www.youtube.com/watch?v=IzWxgFcJFdk&list=PL1tk5lGm7zvQOXi24s586pwDF_yseZ-80&index=7)

[https://www.youtube.com/watch?v=2tCVOkbediU&list=PL1tk5lGm7zvQOXi24s586pwDF\\_yseZ-80&index=9](https://www.youtube.com/watch?v=2tCVOkbediU&list=PL1tk5lGm7zvQOXi24s586pwDF_yseZ-80&index=9)

GT repo: <https://github.com/GridTools>

GT4Py repo: <https://github.com/GridTools/gt4py>